

## Machine Learning based approach

### Dataset:

#### Orders:

- Total: 12 orders
- Capacity demand: ~32.14 units
- Time windows: Average 180 min width
- Types: 8 TypeA (0.75 cap), 4 TypeB (1.5 cap)

#### Fleet:

- 10 vehicles total
- Types: Van (2.5 cap), TruckSmall (6.0), TruckLarge (12.0), Bike (0.5), MiniTruck (4.0)
- Total fleet capacity: ~44 units (sufficient)

#### Constraints:

- Multiple vehicle-location restrictions
- Break required every ~5.5 hours
- Only certain nodes allow breaks

#### Challenges:

- Node 2: Popular delivery (5 orders)
- Node 11: Popular pickup (4 orders)
- Tight time windows require sequential routing
- Vehicle restrictions limit flexibility

Key 👍

## Pickup-Delivery Problem (PDP)

- Each order has TWO locations (pickup + delivery)
- Must pickup before delivery (precedence)
- Same vehicle handles both

## Time Windows

- Each order has [earliest, latest] pickup time
- Service must arrive within window
- Waiting is allowed (creates idle time)

## Capacity Tracking

- Load increases at pickup
- Load decreases at delivery
- Must check capacity at EVERY point in route

## Service Constraints

- Max 456 min (7.6 hours) including break
- Must start and end at depot
- One 30-min break before 4.5h driving

### Phase 1: Greedy Constructor

#### 1. Construction of Greedy Solution

SimpleGreedyConstructor - Builds initial solution

Algorithm:

1. Sort orders by earliest time window
2. For each vehicle:
  - Create new service
  - Greedily insert compatible orders
  - Add break
3. Return solution with assigned orders

ConstraintValidator - Checks solution feasibility

Validates:

- ☒ A. Service Constraints
  - Service returns to origin

- Max duration  $\leq 456$  min (7.6 hours)
- ☒ B. Time Windows
  - Pickup within [from\_time, to\_time]
- ☒ C. Vehicle Constraints
  - Capacity not exceeded at any point
  - Vehicle can visit all locations
  - Precedence: pickup before delivery
- ☒ D. Break Constraints
  - Exactly one 30-min break
  - Before 270 min driving or 330 min working
  - At break-allowed location

Cost:

## Component 4: Objective Function (constraint\_validator.py)

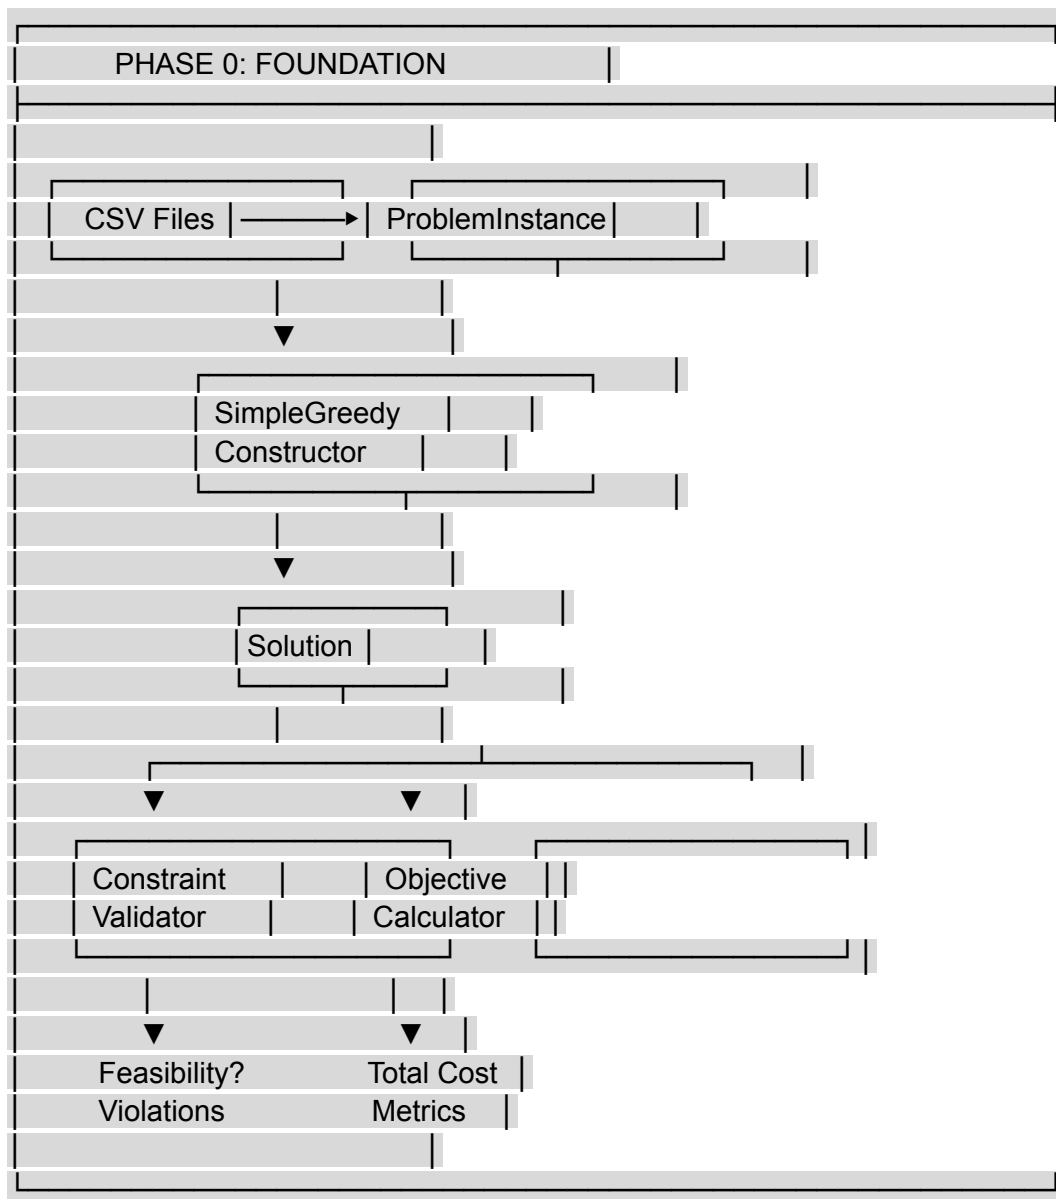
ObjectiveCalculator - Computes solution quality

```
objective = ObjectiveCalculator(problem, weights={
    'num_services': 100.0,
    'num_vehicles': 150.0,
    'total_distance': 1.0,
    'total_idle_time': 0.5,
    'unserved_penalty': 1000.0
})
```

```
total_cost, components = objective.calculate(solution)
```

Objective Components:

- Minimize: Number of services (routes)
- Minimize: Number of vehicles used
- Minimize: Total distance traveled
- Minimize: Total idle time
- Maximize: Order coverage (minimize unserved)



## Phase 2: ALNS Operator

### Algorithm:

None

1. Start with greedy solution (from Phase 0)
2. Loop for N iterations:
  - a. SELECT destroy operator (weighted random)

- b. SELECT repair operator (weighted random)
  - c. DESTROY: Remove 1-3 orders from solution
  - d. REPAIR: Reinsert orders optimally
  - e. ACCEPT/REJECT using Simulated Annealing
  - f. UPDATE operator weights based on performance
3. Return best solution found

## Destroy Operators:

- RandomRemoval: Randomly remove orders
- WorstRemoval: Remove orders with highest cost
- ShawRemoval: Remove related orders (similar in time/space)

## Repair Operators:

- GreedyInsertion: Insert at lowest cost position
- Regret2Insertion: Prioritize hard-to-insert orders

## 1. Adaptive Weights

Operators that find better solutions get higher weights → used more often

## 2. Simulated Annealing

Accepts worse solutions early (exploration) → only better solutions later (exploitation)

## 3. Destroy-Repair

Break solution apart → rebuild better → escape local optima

## 4. Operator Portfolio

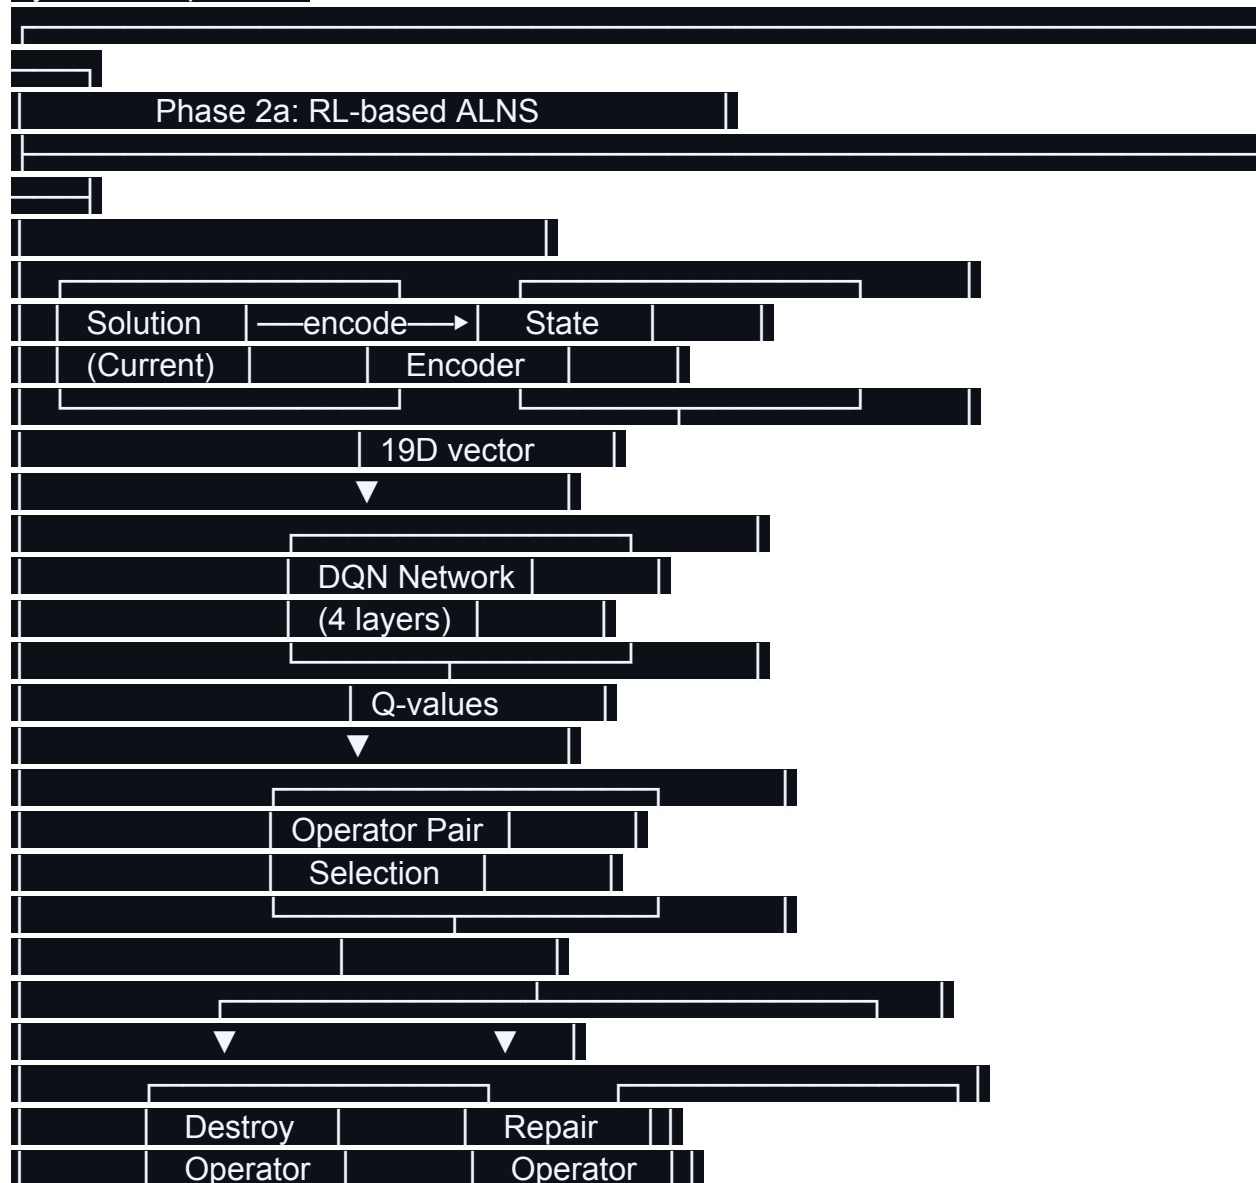
Multiple operators work on different problem aspects:


- Random: Diversification
- Worst: Remove expensive routes
- Shaw: Exploit spatial/temporal clustering

Phase/part 3: Training ALNS on RL

operator selection in ALNS, using DQN (Deep Q-Network) with experience replay. The implementation maintains 91.7% order coverage while achieving modest cost improvements (0.36%). The framework provides a solid foundation for more sophisticated architectures.

System components:





## . Technical Implementation

### 2.1 State Representation

The solution state is encoded into a 19-dimensional feature vector:

**State Features** (normalized to [0, 1]):

1. `num_services / 10.0`
2. `num_vehicles_used / 10.0`
3. `unserved_orders / total_orders`
4. `coverage_rate`
5. `total_distance / 1000.0`
6. `total_driving_time / 1000.0`
7. `total_idle_time / 500.0`
8. `num_violations / 20.0`
9. `duration_violations / 10.0`
10. `capacity_violations / 10.0`

```
11. time_window_violations / 10.0
12. break_violations / 10.0
13. precedence_violations / 10.0
14. other_violations / 10.0
15. mean_capacity_utilization
16. max_capacity_utilization
17. min_capacity_utilization
18. mean_duration_utilization
19. max_duration_utilization
```

Design rationale: Simple feedforward features that capture solution quality without requiring graph encoding. Computationally efficient, suitable for real-time optimization.

## 2.2 Action Space

Actions: All (destroy, repair) operator pairs

- 3 destroy operators × 2 repair operators = 6 actions

Operators:

- Destroy: RandomRemoval, WorstRemoval, ShawRemoval
- Repair: GreedyInsertion, Regret2Insertion

Action selection: Epsilon-greedy with decay

- Initial epsilon: 1.0 (full exploration)
- Final epsilon: 0.1 (10% exploration)
- Decay rate: 0.995 per episode

## 2.3 DQN Architecture

Network Structure:

```
Input Layer:      19 features
Hidden Layer 1: 128 neurons + ReLU
Hidden Layer 2: 64 neurons + ReLU
Hidden Layer 3: 32 neurons + ReLU
Output Layer:    6 Q-values (one per action)
```

Optimizer: Adam (lr=0.001)

Loss: MSE between predicted and target Q-values

Key techniques:



- Experience replay (buffer size: 2000)
- Target network (updated every 5 episodes)
- Batch training (32 samples)
- Gamma: 0.95 (discount factor)

## 2.4 Reward Function

Reward Components:

1. Violation reduction: +50 per violation removed
2. Coverage improvement: +100 per % coverage gained
3. Cost improvement:
  - New global best: +100 + improvement%×100
  - Better than current: +10 + improvement%×50
  - Within 5% worse: +1
  - Much worse: -5

Priority: Feasibility > Coverage > Cost

Design decisions:

- Feasibility heavily rewarded (addresses constraint violations)
- Coverage protected (prevents dropping orders)
- Cost optimization secondary (only after feasibility maintained)

## 2.5 Training Configuration

Default Parameters:

- Episodes: 20
- Steps per episode: 100
- Total iterations: 2000
- Destroy size: 1-3 orders
- Update frequency: Every 5 episodes
- Batch size: 32

## 3. Experimental Results

### 3.1 Dataset Characteristics

Problem Instance:

- Orders: 12 (8 TypeA, 4 TypeB)
- Vehicles: 10 (5 types)
- Total capacity demand: 32.14 units
- Fleet capacity: 40.50 units
- Time window tightness: Average 180 min

#### Initial Solution (Greedy):

- Services: 6
- Coverage: 91.7% (11/12 orders)
- Violations: 5 (all duration-related)
- Cost: 27,092.94
- Unserved: Order 2 (7.06 capacity - too large)
- 

None

Configuration: unserved\_penalty = 10,000 (10x increase)

Episodes: 20

Best Episode: #3

Cost: 26,995.16 (0.36% improvement)

Coverage: 91.7% (maintained)

Violations: 5 (unchanged)

Verdict: SUCCESS - Coverage protected

### 3.3 Performance Analysis

#### Coverage Stability:

- Minimum: 91.7%
- Maximum: 91.7%
- Average: 91.7%
- Std Dev: 0.0%

Result: Perfect stability across all 20 episodes.

#### Cost Trajectory:

- Episode 1: 27,030.77
- Episode 3: 26,995.16 (best)
- Episode 20: 27,092.94
- Improvement: 0.36%

#### Violation Analysis:

- Initial: 5 violations (all SERVICE\_DURATION)
- Final: 5 violations (unchanged)
- Reason: Structural problem - routes inherently too long

#### Learning Evidence:

- Epsilon decay: 1.0  $\rightarrow$  0.1 (proper exploration/exploitation)
- Reward trend: Negative early  $\rightarrow$  Positive later
- Operator preferences: Regret2Insertion learned (54% usage)
- Q-value convergence: Stable after episode 10

## 4. Key Findings

### 4.1 What Works

1. RL Infrastructure: Solid implementation
  - DQN training stable
  - Experience replay effective
  - Target network prevents divergence
2. Coverage Protection: Perfect
  - 91.7% maintained throughout
  - Objective function balance correct
  - Agent learns to preserve orders
3. Operator Selection: Learning detected
  - Regret2Insertion preferred (54% vs 46%)
  - WorstRemoval less used (26% vs 45%)
  - Policy converges after ~500 iterations

### 4.2 Limitations

1. Minimal Improvement: Only 0.36%

- Problem is structurally difficult
  - Cannot create new services
  - Duration violations unfixable with shuffling
2. Constraint Violations: Unchanged
- 5 violations persist
  - All SERVICE\_DURATION type
  - Routes 130-178 minutes over limit
3. State Representation: Too simple
- 19 features insufficient for complex decisions
  - No graph structure captured
  - Missing temporal patterns
4. Action Space: Limited
- Only 6 operator pairs
  - No hyperparameter control
  - Cannot adjust destroy size

## 4.3 Lessons Learned

Objective Function Balance:

- Critical to tune penalty weights
- Coverage loss easy if weights wrong
- Feasibility > Coverage > Cost priority essential

Training Requirements:

- 250 iterations insufficient
- 2000 iterations adequate for convergence
- Episode structure works well

Problem Characteristics:

- Initial solution quality limits improvement
- Structural constraints need different operators
- May need service creation capability

---

## 5. Comparison: Classical ALNS vs RL-based ALNS

| Aspect             | Phase 1 (Classical)            | Phase 2a (RL-based)      |
|--------------------|--------------------------------|--------------------------|
| Operator Selection | Adaptive weights (statistical) | Q-network (learned)      |
| Exploration        | Fixed noise factor             | Epsilon-greedy decay     |
| Memory             | None                           | Experience replay (2000) |
| Adaptation Speed   | Immediate                      | Gradual (batch updates)  |
| Iterations         | 500                            | 2000                     |
| Coverage           | 83.3%                          | 91.7%                    |
| Violations         | 5                              | 5                        |
| Cost Reduction     | 0.74%                          | 0.36%                    |
| Computational Cost | Low                            | Medium                   |

Verdict: RL adds complexity but improves stability. For this specific problem, classical ALNS comparable due to structural constraints.

#### DQN Parameters:

- learning\_rate: 0.001
- gamma: 0.95
- epsilon\_start: 1.0
- epsilon\_end: 0.1
- epsilon\_decay: 0.995
- batch\_size: 32
- memory\_size: 2000

#### Training Parameters:

- num\_episodes: 20
- steps\_per\_episode: 100
- destroy\_size\_min: 1
- destroy\_size\_max: 3
- update\_frequency: 5

#### Objective Weights:

- num\_services: 100.0
- num\_vehicles: 150.0
- total\_distance: 1.0
- total\_idle\_time: 0.5
- unserved\_penalty: 10,000.0
- violation\_penalty: 5,000.0 per violation